

STATS 232A Project: Generator

1 Generator: real inference

The model has the following form:

$$Y = f(Z; W) + \epsilon \quad (1)$$

$$Z \sim N(0, I_d), \epsilon \sim N(0, \sigma^2 I_D), d < D. \quad (2)$$

$f(Z; W)$ maps latent factors into image Y , where W collects all the connection weights and bias terms of the ConvNet.

Adopting the language of the EM algorithm, the complete data model is given by

$$\log p(Y, Z; W) = \log[p(Z)p(Y|Z, W)] \quad (3)$$

$$= -\frac{1}{2\sigma^2} \|Y - f(Z; W)\|^2 - \frac{1}{2} \|Z\|^2 + \text{const}. \quad (4)$$

The observed-data model is obtained by integrating out Z : $p(Y; W) = \int p(Z)p(Y|Z, W)dZ$. The posterior distribution of Z is given by $p(Z|Y, W) = p(Y, Z; W)/p(Y; W) \propto p(Z)p(Y|Z, W)$ as a function of Z .

We want to minimize the observed-data log-likelihood, which is $L(W) = \sum_{i=1}^n \log p(Y_i; W) = \sum_{i=1}^n \log \int p(Y_i, Z_i; W)dZ_i$. The gradient of $L(W)$ can be calculated according to the following well-known fact that underlies the EM algorithm:

$$\frac{\partial}{\partial W} \log p(Y; W) = \frac{1}{P(Y; W)} \frac{\partial}{\partial W} \int p(Y, Z; W)dZ \quad (5)$$

$$= E_{p(Z|Y, W)} \left[\frac{\partial}{\partial W} \log p(Y, Z; W) \right]. \quad (6)$$

The expectation with respect to $p(Z|Y, W)$ can be approximated by drawing samples from $p(Z|Y, W)$ and then compute the Monte Carlo average.

The Langevin dynamics for sampling $Z \sim p(Z|Y, W)$ iterates

$$Z_{\tau+1} = Z_\tau + \delta U_\tau + \frac{\delta^2}{2} \left[\frac{1}{\sigma^2} (Y - f(Z_\tau; W)) \frac{\partial}{\partial Z} f(Z_\tau; W) - Z_\tau \right], \quad (7)$$

where τ denotes the time step for the Langevin sampling, δ is the step size, and U_τ denotes a random vector that follows $N(0, I_d)$.

The stochastic gradient algorithm can be used for learning, where in each iteration, for each Z_i , only a single copy of Z_i is sampled from $p(Z_i|Y_i, W)$ by running a finite number of steps of Langevin dynamics starting from the current value of Z_i , i.e., the warm start. With $\{Z_i\}$ sampled in this manner, we can update the parameter W based on the gradient $L'(W)$, whose Monte Carlo approximation is:

$$L'(W) \approx \sum_{i=1}^n \frac{\partial}{\partial W} \log p(Y_i, Z_i; W) \quad (8)$$

$$= - \sum_{i=1}^n \frac{\partial}{\partial W} \frac{1}{2\sigma^2} \|Y_i - f(Z_i; W)\|^2 \quad (9)$$

$$= \sum_{i=1}^n \frac{1}{\sigma^2} (Y_i - f(Z_i; W)) \frac{\partial}{\partial W} f(Z_i; W). \quad (10)$$

Algorithm 1 describes the details of the learning and sampling algorithm.

Algorithm 1 Generator: real inference

Input:

- (1) training examples $\{Y_i, i = 1, \dots, n\}$,
- (2) number of Langevin steps l ,
- (3) number of learning iterations T .

Output:

- (1) learned parameters W ,
- (2) inferred latent factors $\{Z_i, i = 1, \dots, n\}$.

- 1: Let $t \leftarrow 0$, initialize W .
 - 2: Initialize Z_i , for $i = 1, \dots, n$.
 - 3: **repeat**
 - 4: **Inference step:** For each i , run l steps of Langevin dynamics to sample $Z_i \sim p(Z_i|Y_i, W)$ with warm start, i.e., starting from the current Z_i , each step follows equation 7.
 - 5: **Learning step:** Update $W \leftarrow W + \gamma_t L'(W)$, where $L'(W)$ is computed according to equation 10, with learning rate γ_t .
 - 6: Let $t \leftarrow t + 1$.
 - 7: **until** $t = T$
-

1.1 TO DO

For the lion-tiger category, learn a model with 2-dim latent factor vector. Fill the blank part of `./GenNet/GenNet.py`. **Show:**

- (1) Reconstructed images of training images, using the inferred z from training images.

- (2) Randomly generated images, using randomly sampled z .
- (3) Generated images with linearly interpolated latent factors from $(-2, 2)$ to $(-2, 2)$.
For example, you interpolate 8 points from $(-2, 2)$ for each dimension of z . Then you will get a 8×8 panel of images. You should be able to see that tigers slight change to lion.
- (4) Plot of loss over iteration.

What to submit

Write a report to show your results. And zip the report with your code.